


Big Data Ecosystem



5. NoSQL with HBase

What is NoSQL?

- **Key-value:**

- 1 key = 1 value
- Examples : Redis, Memcache (in memory)

- **Columnar:**

- 1 key = X values stored in columns
- Examples : HBase, Cassandra

What is NoSQL?

- **Document stores:**

- JSON/XML like objects
- Examples : MongoDB, CouchDB

- **Graph:**

- Store nodes and relationships between them
- Examples : Neo4j, JanusGraph

Apache HBase

- *Apache HBase™ is the Hadoop database, a distributed, scalable, big data store - hbase.apache.org*
- Random, realtime read/write access to Big Data
- Inspired from Google BigTable paper (2006)
- Stores data in HDFS

The CAP Theorem

- **Consistency:** For one given query, all the nodes return the same result
- **Availability:** Every query receives a fast response, without guarantee that this is the latest value
- **Partition tolerance:** The system continues to work even if some nodes are disconnected
- No distributed storage system can ensure more than 2 of those properties at the same time

Apache HBase

- HBase is a **CP columnar** database
- AP equivalent = [Apache Cassandra](#)

HBase: Data structure

- A table is a collection of rows
- A row has several **column families** defined at table creation
- A column family has **any number of columns**
- 1 row = { **key**: {
 column_family_1: { col_1: a, col_2: b},
 column_family_2: { col_3: c, col_4: d}
}}

HBase: Data storage

- Data is stored:
 - In HDFS in the **HFiles**
 - In RAM in the **Memstore**
- Tables are split into **regions**
 - 1 region = all rows in a certain **range of keys**
 - The number of regions depends on the size of the table

HBase: Components

- **HBase Master**

- Handle table creation/deletion queries
- **Assigns regions** to RegionServers and monitor them

- **RegionServer**

- 1 region of a table is managed by 1 RegionServer

HBase: Data storage in RegionServers

1. On data write:
 - a. Storing in the **Memstore in RAM**
 - b. Storing in the **HLog in HDFS** (Write Ahead Log)
2. When the Memstore or the HLog reaches max size: flush to **HFile in HDFS** (1 per CF). Temporary **compactions**
3. If the RegionServer crashes, operations are replayed from the HLog for the Memstore and metadata loaded from the HFile (by another RegionServer)

HBase: Partition tolerance and HA

- RegionServers tracked by ZooKeeper
 - On crash, HMaster knows it through ZK
 - HMaster re-assigns regions to other RegionServers
- High Availability: multiple HBase Masters

Querying HBase

- HBase client
 - CLI
 - Languages / tools (e.g. NiFi)
- Apache Phoenix, SQL on HBase (in memory)
- Apache Hive